

---

# Zeek

## Network Security Monitor

Maruthi S. Inukonda

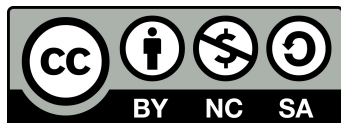
08<sup>th</sup> March 2024

---

# License

Copyright © 2024 - Maruthi S. Inukonda

This work is licensed under a Creative Common Attribution-NonCommercial-ShareAlike 4.0 Unported License. This license is available at <https://creativecommons.org/licenses/by-nc-sa/4.0/>



# Agenda

- Introduction
- Installation
- Logs
- Scripting

# Introduction

# Zeek

- Zeek is a Network Security Monitoring Tool
- Formerly known as Bro
- Developed at Lawrence Berkely National Laborator (LBNL), USA
- Published in Unix Security Symposium 1998



USENIX

<https://www.usenix.org> > sec98 > paxson > paxson PDF



## Bro: A System for Detecting Network Intruders in Real-Time

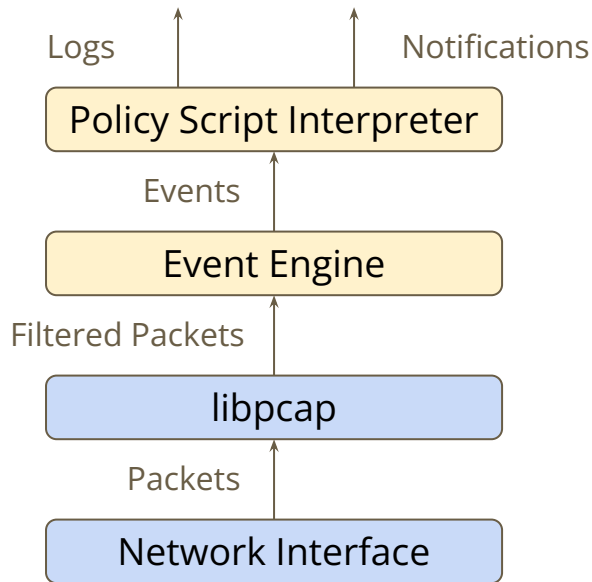
by V Paxson · Cited by **4047** — We describe **Bro**, a stand-alone system for detecting network intruders in real-time by passively monitoring a network link over which the intruder's ...

22 pages

- Code open-sourced under BSD License
- Not an active security device like Firewall, Intrusion Prevention System
- Unobtrusively observes network traffic
- Used for in products for Evidence-Based Network Detection and Response (NDR), and Threat Hunting, Log Management (see Corelight, Security Onion)

# Architecture

- **Network Interface**
  - Tap network link passively, send up a copy of all traffic
- **Libpcap**
  - Kernel filters down high-volume stream via packet capture library
- **Event Engine**
  - Distills filtered stream of packets into high-level events
  - Developed in C++
- **Policy Script Interpreter**
  - Loads scripts written in domain specific policy language (similar to Python)
  - Runs event handlers corresponding to events to take action



# Installation

## Setup Apt Repository & Install

- Install Long Term Support (LTS) version  
<https://docs.zeek.org/en/lts/install.html>
- Packages are available for major distros of Linux
- You may also build from source.

```
$ sudo apt-get install -y --no-install-recommends g++ cmake make libpcap-dev
```

```
$ echo 'deb http://download.opensuse.org/repositories/security:zeek/xUbuntu_20.04/' | sudo tee /etc/apt/sources.list.d/security:zeek.list
```

```
$ curl -fsSL  
https://download.opensuse.org/repositories/security:zeek/xUbuntu_20.04/Release.key |  
gpg --dearmor | sudo tee /etc/apt/trusted.gpg.d/security_zeek.gpg > /dev/null
```

```
$ sudo apt update
```

```
$ sudo apt install zeek-lts
```

Select appropriate mail server. If you don't have one, select "No configuration"



## Configuration (1/2)

- Zeek runs as root user. Switch to root shell
- Configure root's shell environment and reload it
- Verify the shell environment

```
$ sudo su -
```

```
#
```

```
# vi /root/.bashrc
```

```
.
```

```
.
```

```
.
```

```
export PATH=$PATH:/opt/zeek/bin/
```

```
# exec bash
```

```
# which zeek
```

```
/opt/zeek/bin/zeek
```

```
# which zeekctl
```

```
/opt/zeek/bin/zeekctl
```

```
# which zeek-cut
```

```
/opt/zeek/bin/zeek-cut
```

## Configuration (2/2)

- Find your network interface

```
# ip link show
```

```
.  
2: enp0s7: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP mode  
DEFAULT group default qlen 1000  
    link/ether be:ef:fe:ed:be:ef brd ff:ff:ff:ff:ff:ff  
3: wlp0s20f3: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP mode  
DORMANT group default qlen 1000  
    link/ether de:ad:be:ef:de:ad brd ff:ff:ff:ff:ff:ff
```

- Find your network interface

```
# vi /opt/zeek/etc/node.cfg
```

```
.  
.  
[zeek]  
type=standalone  
host=localhost  
interface=af_packet::enp0s7
```

## Deploy Zeek

- Deploy Zeek in the background

```
root@terak:~# zeekctl deploy
checking configurations ...
installing ...
removing old policies in
/opt/zeek/spool/installed-scripts-do-not-touch/site ...
removing old policies in
/opt/zeek/spool/installed-scripts-do-not-touch/auto ...
creating policy directories ...
installing site policies ...
generating standalone-layout.zeek ...
generating local-networks.zeek ...
generating zeekctl-config.zeek ...
generating zeekctl-config.sh ...
stopping ...
stopping zeek ...
starting ...
starting zeek ...
```

## Verify Zeek

- Check status of zeek running in the background

```
# zeekctl status
```

```
waiting for lock (owned by PID 24496) ...
```

Name	Type	Host	Status	Pid	Started
zeek	standalone	localhost	running	25670	04 Mar 22:37:35

- Check network statistics

```
# zeekctl netstats
```

```
zeek: 1709573082.768830 recvd=9714926 dropped=0 link=9715276
```

- Check capture statistics

```
# zeekctl capstats
```

Interface	kpps	mbps	(10s average)
localhost/af_packet::enp0s7	86.2		967.1

# Logs

# Logs

- Currently generated logs are stored in current directory
- Periodically log files are archived in directory named with <date>

```
# ls /opt/zeek/logs/current
```

```
capture_loss.log conn.log dhcp.log  
dns.log files.log http.log ntp.log ocsf.log  
software.log ssl.log stats.log stderr.log stdout.log  
telemetry.log weird.log x509.log
```

```
# ls /opt/zeek/logs
```

```
2024-03-04 current
```

# Logs

- All logs are space/tab separated columns
- For eg., conn.log contains flow records

```
# less /opt/zeek/logs/current/conn.log
```

```
•
•
#fields      ts      uid      id.orig_h  id.orig_p  id.resp_h  id.resp_p  proto  service
duration    orig_bytes  resp_bytes  conn_state  local_orig  local_resp
missed_bytes  history    orig_pkts  orig_ip_bytes  resp_pkts  resp_ip_bytes
tunnel_parents
#types      time      string    addr      port      addr      port      enum    string    interval
count      count    string    bool      bool      count     string    count    count    count
count      set[string]
1 1709602202.400788      CPjdaz38Qfyz0sX9c      172.16.0.228      5353      224.0.0.251      5353      udp
dns      2.010773      330      0      S0      T      F      0      D      3      414      0      0-
2 1709602182.244735      CPR4Mc1NN1j4IBDD72      172.16.0.237      34492      142.250.193.170      443
udp      -      0.428834      3546      4986      SF      T      F      0      Dd      5      3686      9      5238
-
•
•
```

Difficult  
to read

## Extracting Columns

- zeek-cut is the tool to pick required columns

Using input redirection

```
# zeek-cut -m -d ts uid id.orig_h id.orig_p id.resp_h id.resp_p proto < conn.log
ts      uid      id.orig_h  id.orig_p  id.resp_h  id.resp_p  proto
2024-03-05T07:00:02+0530  CPjdaz38Qfyz0sX9c  172.16.0.228  5353  224.0.0.251  5353  udp
2024-03-05T06:59:42+0530  CPR4Mc1NN1j4IBDD72  172.16.0.237  34492  142.250.193.170  443  udp
.
```

Using pipe

```
# cat conn.log | zeek-cut -m -d ts id.orig_h id.orig_p id.resp_h id.resp_p proto
ts      uid      id.orig_h  id.orig_p  id.resp_h  id.resp_p  proto
2024-03-05T07:00:02+0530  CPjdaz38Qfyz0sX9c  172.16.0.228  5353  224.0.0.251  5353  udp
2024-03-05T06:59:42+0530  CPR4Mc1NN1j4IBDD72  172.16.0.237  34492  142.250.193.170  443  udp
.
```



# Scripting

## Script init and done (1/4)

- Create a directory to place all our scripts
- `zeek_init()` event handler is used for initialization
- It is invoked
  - During “`zeekctl deploy`”
  - During “`zeekctl start`”
  - Beginning of “`zeek -r test.pcap`”
- `zeek_done()` is used for cleanup
- It is invoked
  - During “`zeekctl stop`”
  - Ending of “`zeek -r test.pcap`”

```
# mkdir /opt/zeek/share/zeek/site/myscripts
# cd /opt/zeek/share/zeek/site/myscripts

# vi sample.zeek

event zeek_init() {
    print "Zeek started!";
    # Any other initialization stuff goes here
}

event zeek_done() {
    # Any other cleanup stuff goes here
    print fmt("Zeek stopped!");
}
```

## Script init and done (2/4)

- Create a dummy pcap file      `# tcpdump -w /tmp/test.pcap`
- Analyze a pcap file      `# zeek -C -r /tmp/test.pcap sample.zeek`  
Zeek started!  
Zeek stopped!

## Script init and done (3/4)

- By default no logs are printed during daemon/background mode

- By default the global variable `Log::print_to_log` is set to `Log::REDIRECT_NONE`

- To send all prints to `print.log`, edit our script and assign it with `Log::REDIRECT_ALL`

```
# vi sample.zeek
```

```
redef Log::print_to_log = Log::REDIRECT_ALL;
```

```
event zeek_init() {  
    print "Zeek started!";  
    # Any other initialization stuff goes here  
}
```

```
event zeek_done() {  
    # Any other cleanup stuff goes here  
    print fmt("Zeek stopped!");  
}
```

## Script init and done (4/4)

- Add the following line at the **end of the default** site/local.zeeb file

```
# vi /opt/zeek/share/zeek/site/local.zeeb
.  
.  
.  
@load myscripts/sample
```

- Deploy the script
- Check the prints in print.log

```
# zeekctl deploy
```

```
# less /opt/zeek/logs/current/print.log
```

# Variables and Constants

- Global variables are accessible across event handlers and scripts
- Local variables are defined and accessible only inside an event handler
- Constants are accessible in their defined scope (global or local) and cannot be modified

```
# vi sample.zEEK
```

```
redef Log::print_to_log = Log::REDIRECT_ALL;
```

```
global gs = "apple";  
const cs = "mango";
```

```
event zeek_init() {  
    print "Zeek started!";  
    local ls = "orange";  
    print fmt("gs:%s cs:%s ls:%s", gs, cs, ls);  
    gs = "muskmelon";  
    ls = "banana";  
    print fmt("gs:%s ls:%s", gs, ls);  
}
```

```
event zeek_done() {  
    print fmt("gs:%s cs:%s", gs, cs);  
    print fmt("Zeek stopped!");  
}
```

## Primitive Data Types (1/2)

- String
- Boolean
- Integer (64-bit unsigned)
- Double (64-bit double precision)
- Counter
- Timestamp
- Interval

```
local ls: string = "pomegranate";
```

```
local lb: bool = F;
```

```
local li: int = 4;
```

```
global ld: double = 9.25;
```

```
global gc: count = 0;
```

```
local lt: ts = current_time();
```

```
local linvl: ts = 2.5sec;
```

## Primitive Data Types (2/2)

- IP Address
- Subnet
- Port

```
event zeek_init()
{
    .
    .
    .

    local ip : addr = 192.168.1.100;
    local sn : subnet = 192.168.0.0/16;
    local prt : port = 53/udp;
    print(fmt("ip:%s sn:%s port:%d", ip, sn, prt));
}
```



## Control Statements (1/2)

- If - else statement

```
local ip: addr = 192.168.1.100;
local sn: subnet = 192.168.0.0/16;
if ( ip/16 == sn )
    print(fmt("ip:%s belongs to sn %s", ip, sn));
```

- for loop  
No need to define loop  
control variable

```
for (c in "Hello") {
    print fmt("c: %s", c);
}
```

- while loop

```
local i = 0;
while (i <= 5) {
    ++i;
    print "i: ", i;
}
```

## Control Statements (2/2)

- switch case statement

```
local p : int = 3
switch (i) {
  case 1:
    print "i is 1";
    break;
  case 2:
    print "i is 2";
    break;
  case 3:
    print "i is 3";
    break;
  default:
    print "i is not 1-3";
    break;
}
```

## Composite Data Types (1/3)

- Record types are used to store related information together

```
type MyRecordType: record {  
    c: count;  
    s: string;  
};  
  
event zeek_init()  
{  
    .  
    .  
    .  
    local lr : MyRecordType = record($c = 9, $s =  
"nine");  
    print(fmt("lr$c:%d lr$s:%s", lr$c, lr$s));  
    .  
    .  
}
```

## Composite Data Types (2/3)

- Vector

```
local vec = vector("one", "two", "three");
print(fmt("vec: %s %s %s", vec[0], vec[1], vec[2]));
for (v in vec){
    print fmt("v: %s", vec[v]);
}
```

- Set

```
local s: set[port] = { 21/tcp, 23/tcp, 80/tcp,
443/tcp };
for (s in s1){
    print fmt("s: %s", s);
}
if (21/tcp in s1) {
    print("port no exists in set");
}
```

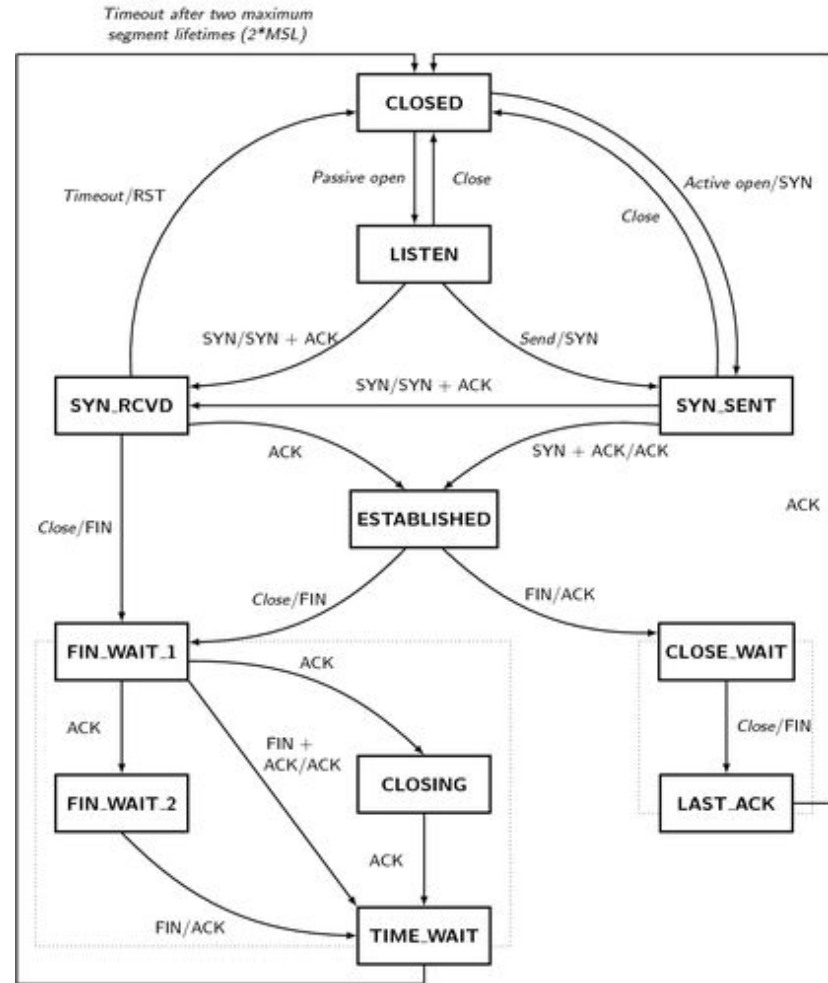
## Composite Data Types (3/3)

- Table

```
local t1: table[count] of string = {  
    [11] = "eleven",  
    [5] = "five",  
};  
for (t in t1){  
    print fmt("t: %s", t);  
}  
if (11 in t1) {  
    print("count exists in table");  
}
```

# TCP State machine

- `connection_established`: Upon receiving the ACK, both the client and server transition into the **ESTABLISHED** state, signifying a successful connection establishment.
  - Applicable for TCP only
- `connection_finished`: Upon receiving the FIN, both the client and server transition into the **CLOSING** state, signifying a successful connection termination.
  - Applicable for TCP only
- `new_connection`: Upon receiving a packet that makes zeek create a new flow tracking data structure.
  - Applicable for both TCP and UDP



# Connection Events

- Add these handlers to the sample.zeeK script

```
# vi sample.zeeK
.
# Applicable for TCP
event connection_established(c: connection)
{
    print fmt("Zeek connection established from %s:%s to %s:%s",
              c$id$orig_h, c$id$orig_p, c$id$resp_h, c$id$resp_p);
}

# Applicable for TCP
event connection_finished(c: connection)
{
    print fmt("Zeek connection finished from %s:%s to %s:%s",
              c$id$orig_h, c$id$orig_p, c$id$resp_h, c$id$resp_p);
}

# Applicable for TCP and UDP
event new_connection(c: connection)
{
    print fmt("Zeek new connection from %s:%s to %s:%s",
              c$id$orig_h, c$id$orig_p, c$id$resp_h, c$id$resp_p);
}
```

## Connection Events

- Run the script on test pcap file

```
# zeek -C -r /tmp/test.pcap sample.zeek
```

```
Zeek new connection from 10.10.20.11:34206/tcp to  
192.168.137.179:22/tcp
```

```
Zeek connection established from 10.10.20.11:56474/tcp to  
192.168.137.179:22/tcp
```

```
Zeek connection finished from 10.10.20.11:56474/tcp to  
192.168.137.179:22/tcp
```



## Packet Events

- `tcp_packet` : Generated for every TCP packet.
- `udp_request`: Generated for each packet sent by a UDP flow's originator.
- `udp_reply`: Generated for each packet sent by a UDP flow's responder.
- A very low-level and expensive event due to the volume of traffic.
- Should be avoided in background mode.

```
# vi sample.zeek
```

```
:
:
event tcp_packet(c: connection, is_orig: bool, flags: string,
seq: count, ack: count, len: count, payload: string)
{
    print fmt("Zeek new tcp packet from %s:%s to %s:%s",
              c$id$orig_h, c$id$orig_p,
              c$id$resp_h, c$id$resp_p);
}

event udp_request(u: connection)
{
    print fmt("Zeek new udp request from %s:%s to %s:%s",
              u$id$orig_h, u$id$orig_p,
              u$id$resp_h, u$id$resp_p);
}

event udp_reply(u: connection)
{
    print fmt("Zeek new udp reply from %s:%s to %s:%s",
              u$id$orig_h, u$id$orig_p,
              u$id$resp_h, u$id$resp_p);
}
```

## Packet Events

- Count the number of TCP/UDP events using zeek
- Count number of TCP/UDP packets using tcpdump utility.

```
# tcpdump -r /tmp/test.pcap udp or tcp | wc -l
reading from file /tmp/test.pcap, link-type EN10MB (Ethernet),
snapshot length 262144
24
```

```
# zeek -C -r /tmp/test.pcap sample.zeek | wc -l
24
```

## Application Protocol-level Events

- Eg. Count the number of http errors

```
# vi sample.zeek
```

```
global http_404_count: count = 0;
```

```
.
```

```
event zeek_done() {
```

```
    print fmt("Zeek stopped. Count: %d", http_404_count);
```

```
}
```

- Redeploy zeek

```
# Applicable for http reply
```

```
event http_reply(c: connection, version: string, code: count,  
    reason: string) {
```

```
    if (code == 404) {
```

```
        # Increment the count
```

```
        http_404_count += 1;
```

- Open a non-existent web page

```
        # Print information or perform other actions
```

```
        print fmt("%d HTTP 404 responses detected.",
```

```
            http_404_count);
```

```
    }
```

- Check the print.log

```
}
```

# References

# Installation

- Book of Zeek (LTS : v6.0.3)  
<https://docs.zeek.org/en/lts/>
- Zeek Training Materials/Products  
<https://github.com/zeek/zeek-training>
- Zeek in Action - Youtube Channel - Many playlists  
<https://www.youtube.com/@Zeekurity>

# Q & A